

Azure Databricks Coding Challenge

Submitted By: Aniroop Gupta
DE Batch 1

1. Creation of the Cluster:

Created a Databricks cluster with appropriate configurations to execute the notebook efficiently. Attached the notebook to the cluster for seamless code execution.

To create a cluster:

- Log in to Azure Databricks.
- Go to the Compute section and click on Create Cluster.
 - Cluster Name: azuser2360_mml.local's Cluster
 - Databricks Runtime Version: 15.4 LTS(Scala 2.12, Spark 3.5.0)
 - Node Type: Single Node, Standard_D4ads_v5
- Click Create Cluster and wait until the status is "Running."

The screenshot shows the 'New compute' page in the Azure Databricks interface. The cluster name is 'azuser2360_mml.local's Cluster'. The configuration is set to 'Single node' and 'Single user access'. The 'Access mode' is 'Single user' and the 'User' is 'azuser2360_mml.local'. The 'Performance' section shows 'Databricks runtime version' as 'Runtime: 15.4 LTS (Scala 2.12, Spark 3.5.0)'. The 'Node type' is 'Standard_D4ads_v5' with '16 GB Memory, 4 Cores'. The 'Terminate after' checkbox is checked, set to '60 minutes of inactivity'. The 'Tags' section is empty. A 'Summary' panel on the right shows: 1 Driver, 16 GB Memory, 4 Cores; Runtime: 15.4.x-scala2.12; Standard_D4ads_v5, 1 DBU/h. At the bottom, there are 'Create compute' and 'Cancel' buttons.

Compute > New compute

azuser2360_mml.local's Cluster

☐ Multi node ☒ Single node

Access mode Single user access

Single user azuser2360_mml.local

Performance

Databricks runtime version

Runtime: 15.4 LTS (Scala 2.12, Spark 3.5.0)

☐ Use Photon Acceleration

Node type

Standard_D4ads_v5 16 GB Memory, 4 Cores

☒ Terminate after 60 minutes of inactivity

Tags

Add tags

Key Value Add

> Automatically added tags

Create compute Cancel

Summary

1 Driver	16 GB Memory, 4 Cores
Runtime	15.4.x-scala2.12
Standard_D4ads_v5	1 DBU/h

2. Create a DataFrame from a Databricks Dataset:

Loaded the dataset mudah-apartment-kl-selangor (1).csv into a Spark DataFrame using the `spark.read.csv` function. Displayed the first 5 rows of the DataFrame to ensure successful loading.

```
1 minute ago 1 Python
```

```
# Load Spark and access Databricks dataset
from pyspark.sql import SparkSession

# Initialize SparkSession (usually auto-initialized in Databricks)
spark = SparkSession.builder.getOrCreate()

# Load sample data
filepath = "/FileStore/tables/mudah_apartment_kl_selangor__1_-1.csv"
df = spark.read.csv(filepath, header=True, inferSchema=True)

# Show the first few rows
df.show(5)
```

▶ (3) Spark Jobs

df: pyspark.sql.dataframe.DataFrame = [ads_id: integer, prop_name: string ... 12 more fields]

ads_id	prop_name	completion_year	monthly_rent	location	property_type	rooms	parking	bathroom	size
100323185	The Hipster @ Tam...	2022.0	RM 4 200 per month	Kuala Lumpur - Ta...	Condominium	5	2.0	6.0	1842 sq.ft.
100203973	Segar Courts	NULL	RM 2 300 per month	Kuala Lumpur - Ch...	Condominium	3	1.0	2.0	1170 sq.ft.
100323128	Pangsapuri Terata...	NULL	RM 1 000 per month	Kuala Lumpur - Ta...	Apartment	3	NULL	2.0	650 sq.ft.
100191767	Sentul Point Suit...	2020.0	RM 1 700 per month	Kuala Lumpur - Se...	Apartment	2	1.0	2.0	743 sq.ft.

only showing top 5 rows

3. View Schema of the Data:

Displayed the schema of the dataset using the `printSchema()` function to understand the column names and data types.

```
Just now (<1s) 2 Python
```

```
# Check the structure and data types of the DataFrame
df.printSchema()
```

```
root
 |-- ads_id: integer (nullable = true)
 |-- prop_name: string (nullable = true)
 |-- completion_year: double (nullable = true)
 |-- monthly_rent: string (nullable = true)
 |-- location: string (nullable = true)
 |-- property_type: string (nullable = true)
 |-- rooms: string (nullable = true)
 |-- parking: double (nullable = true)
 |-- bathroom: double (nullable = true)
 |-- size: string (nullable = true)
 |-- furnished: string (nullable = true)
 |-- facilities: string (nullable = true)
 |-- additional_facilities: string (nullable = true)
 |-- region: string (nullable = true)
```

4. Display Summary Statistics:

Used the describe() function to generate summary statistics for numeric columns, providing insights into the data distribution.

```
Python Just now (4s) 3

# Show summary statistics for numerical columns
df.describe().show()
```

▶ (2) Spark Jobs

summary	ads_id	prop_name	completion_year	monthly_rent	location	property_type	rooms
parking	bathroom	size	furnished	facilities	additional_facilities	region	
count	19991	19043	10806	19989	19991	19991	19985
14289	19985	19991	19986	17782	14043	19991	
mean	9.97067053108899E7	NULL	2014.822320932815	NULL	NULL	NULL	2.680278236501026
8241304499964	1.8917187890918188	NULL	NULL	NULL	NULL	NULL	1.416
stddev	3482574.7035407154	NULL	6.735354696429903	NULL	NULL	NULL	0.8078379940634384
3677630463505	0.5562656902904315	NULL	NULL	NULL	NULL	NULL	0.567
min	16525511	1 Harmonis	1977.0	RM 1 000 per month	Kuala Lumpur - Am...	Apartment	1
1.0	1.0	1 sq.ft.	Fully Furnished	Barbeque area	, Air-Cond, Cooki...	Kuala Lumpur	
max	100854617	wangsa maju secti...	2025.0	RM 999 per month	Selangor - Ulu Klang	Townhouse Condo	More than 10
10.0	8.0	999999999 sq.ft.	Partially Furnished	Tennis Court, Swi...	Washing Machine, ...	Selangor	

5. Count Rows in Dataset:

Counted the total number of rows in the dataset using the count() function to understand the dataset's size.

```
Python Just now (1s) 4

# Count the number of rows in the dataset
row_count = df.count()
print(f"Total rows in the dataset: {row_count}")
```

▶ (2) Spark Jobs

Total rows in the dataset: 19991

6. Group and Aggregate Data Sort Data:

Grouped the dataset by the location column and calculated the average monthly_rent to analyze rent trends.

```
Python Just now (1s) 5

# Group data by 'location' and count the number of listings
location_count_df = df.groupBy("location").count()

# Display the result for visualization
location_count_df.show(10)
```

▶ (2) Spark Jobs

location_count_df: pyspark.sql.dataframe.DataFrame = [location: string, count: long]

location	count
Kuala Lumpur - OUG	63
Kuala Lumpur - Ji...	9
Selangor - Kota D...	172
Selangor - Puncak...	63
Selangor - Bandar...	32
Selangor - Glenmarie	19
Selangor - Pulau ...	1
Kuala Lumpur - So...	60
Kuala Lumpur - Ch...	1
Selangor - Petali...	612

only showing top 10 rows

7. Sort Data:

Sorted the aggregated data in descending order of location count using the `orderBy()` function to identify the count of locations.

```
Just now (1s) 6 Python
```

```
# Sort by count in descending order
sorted_df = location_count_df.orderBy("count", ascending=False)
display(sorted_df.limit(5))
```

(2) Spark Jobs

sorted_df: pyspark.sql.dataframe.DataFrame = [location: string, count: long]

Table +

	location	count
1	Kuala Lumpur - Cheras	1623
2	Selangor - Kajang	1022
3	Kuala Lumpur - Setap...	973
4	Selangor - Shah Alam	971
5	Selangor - Cyberjaya	879

8. Top 10 Most Frequent Monthly Rent Values:

Grouped by `monthly_rent` and counted occurrences, then displayed the top 10 most frequent rent values in descending order.

```
Just now (1s) 8 Python
```

```
# Group by 'monthly_rent' and count occurrences, then sort by count in descending order
top_rent_values = df.groupBy("monthly_rent").count().orderBy("count", ascending=False)

# Show the top 10 most frequent rent values
top_rent_values.show(10)

# Visualize the data
display(top_rent_values)
```

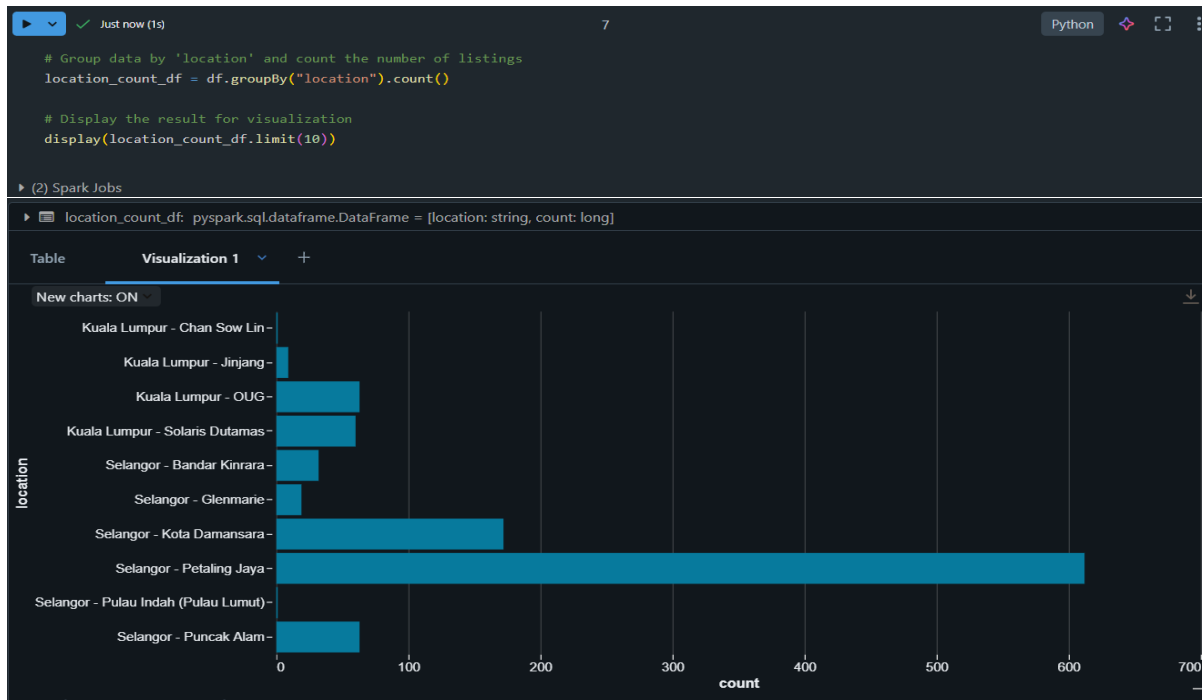
(4) Spark Jobs

top_rent_values: pyspark.sql.dataframe.DataFrame = [monthly_rent: string, count: long]

```
+-----+
|  monthly_rent|count|
+-----+
|RM 1 500 per month| 1401|
|RM 1 200 per month| 1373|
|RM 1 300 per month| 1209|
|RM 1 600 per month| 1067|
|RM 1 400 per month| 1024|
|RM 1 000 per month|  986|
|RM 1 800 per month|  899|
|RM 1 100 per month|  828|
|RM 1 700 per month|  755|
|RM 2 000 per month|  728|
+-----+
only showing top 10 rows
```

9. Create a Visualizations:

Visualized the count of listings by location using a bar chart to analyze the distribution of data across different areas.



10. Another Visualization:

Created a bar chart to display the distribution of monthly rent values, showing the frequency of each rent amount.

