

Azure Databricks Assignment

Submitted By: Aniroop Gupta
DE Batch 1

1. Write Data to CSV, JSON, Parquet, delta:

```
▶ ✓ 05:31 PM (<1s) 1 Python
```

```
# Sample DataFrame
data = [("Alice", 34), ("Bob", 45), ("Cathy", 25)]
columns = ["Name", "Age"]
df = spark.createDataFrame(data, columns)
```

```
▶ df: pyspark.sql.dataframe.DataFrame = [Name: string, Age: long]
```

```
▶ ✓ 05:32 PM (1s) 2
```

```
# 1. Write DataFrame to CSV
df.write.option("header", "true").csv("/dbfs/mnt/path/to/outputs/csv")
```

```
▶ (1) Spark Jobs
```

```
▶ ✓ 05:32 PM (1s) 3
```

```
# 2. Write DataFrame to JSON
df.write.json("/dbfs/mnt/path/to/outputs/json")
```

```
▶ (1) Spark Jobs
```

```
▶ ✓ 05:32 PM (1s) 4
```

```
# 3. Write DataFrame to Parquet
df.write.parquet("/dbfs/mnt/path/to/outputs/parquet")
```

```
▶ (1) Spark Jobs
```

05:33 PM (4s)

5

```
# 4. Write DataFrame to Delta format
df.write.format("delta").save("/dbfs/mnt/path/to/outputs/delta")
```

(6) Spark Jobs

05:34 PM (5s)

6

```
# 5. Optionally create a Delta table
df.write.format("delta").mode("overwrite").saveAsTable("people_deltatable")
```

(6) Spark Jobs

2. Writing dataframe to Delta Table:

05:35 PM (4s)

1

Python



```
# Sample DataFrame
data = [("Alice", 34), ("Bob", 45), ("Cathy", 25)]
columns = ["Name", "Age"]
df = spark.createDataFrame(data, columns)

# Write DataFrame to Delta table
df.write.format("delta").mode("overwrite").save("/dbfs/mnt/path/to/delta_table")
```

(6) Spark Jobs

df: pyspark.sql.dataframe.DataFrame = [Name: string, Age: long]

3. Exploratory data analysis (EDA) in Databricks:

05:43 PM (<1s)

1

Python



```
import pandas as pd

spark_df = spark.read.format("delta").load("/dbfs:/user/hive/warehouse/loan_data")

data = spark_df.toPandas()

print(data)
```

(1) Spark Jobs

spark_df: pyspark.sql.dataframe.DataFrame = [Loan_ID: string, Gender: string ... 11 more fields]

	Loan_ID	Gender	Married	...	Credit_History	Property_Area	Loan_Status
0	LP001002	Male	No	...	1.0	Urban	Y
1	LP001003	Male	Yes	...	1.0	Rural	N
2	LP001005	Male	Yes	...	1.0	Urban	Y
3	LP001006	Male	Yes	...	1.0	Urban	Y
4	LP001008	Male	No	...	1.0	Urban	Y
...
609	LP002978	Female	No	...	1.0	Rural	Y
610	LP002979	Male	Yes	...	1.0	Rural	Y
611	LP002983	Male	Yes	...	1.0	Urban	Y
612	LP002984	Male	Yes	...	1.0	Urban	Y
613	LP002990	Female	No	...	0.0	Semiurban	N

[614 rows x 13 columns]

05:44 PM (<1s)

2

```
#Printing rows of data & display values
display(data.head()) #first 5 rows
display(data.tail()) #last 5 rows
```

Table +



	A ^B _C Loan_ID	A ^B _C Gender	A ^B _C Married	A ^B _C Dependents	A ^B _C Education	A ^B _C Self_Employed	1 ² ₃ ApplicantIncome	1.2 CoapplicantIncome
1	LP001002	Male	No	0	Graduate	No	5849	
2	LP001003	Male	Yes	1	Graduate	No	4583	
3	LP001005	Male	Yes	0	Graduate	Yes	3000	
4	LP001006	Male	Yes	0	Not Graduate	No	2583	

5 rows | 0.22 seconds runtime

Refreshed 40 minutes ago

Table +



	A ^B _C Loan_ID	A ^B _C Gender	A ^B _C Married	A ^B _C Dependents	A ^B _C Education	A ^B _C Self_Employed	1 ² ₃ ApplicantIncome	1.2 CoapplicantIncome
1	LP002978	Female	No	0	Graduate	No	2900	
2	LP002979	Male	Yes	3+	Graduate	No	4106	
3	LP002983	Male	Yes	1	Graduate	No	8072	
4	LP002984	Male	Yes	2	Graduate	No	7583	

05:44 PM (<1s)

3

```
#Printing the column names of the DataFrame
list(data.columns)
```

```
['Loan_ID',
 'Gender',
 'Married',
 'Dependents',
 'Education',
 'Self_Employed',
 'ApplicantIncome',
 'CoapplicantIncome',
 'LoanAmount',
 'Loan_Amount_Term',
 'Credit_History',
 'Property_Area',
 'Loan_Status']
```

05:44 PM (<1s)

5

```
#Descriptive Statistical Measures of a DataFrame
data.describe()
```

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History
count	614.000000	614.000000	592.000000	600.000000	564.000000
mean	5403.459283	1621.245798	146.412162	342.000000	0.842199
std	6109.041673	2926.248369	85.587325	65.12041	0.364878
min	150.000000	0.000000	9.000000	12.000000	0.000000
25%	2877.500000	0.000000	100.000000	360.000000	1.000000
50%	3812.500000	1188.500000	128.000000	360.000000	1.000000
75%	5795.000000	2297.250000	168.000000	360.000000	1.000000
max	81000.000000	41667.000000	700.000000	480.000000	1.000000

```
#Missing Data Handling
data.dropna()
```

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History	Property
1	LP001003	Male	Yes	1	Graduate	No	4583	1508.0	128.0	360.0	1.0	
2	LP001005	Male	Yes	0	Graduate	Yes	3000	0.0	66.0	360.0	1.0	
3	LP001006	Male	Yes	0	Not Graduate	No	2583	2358.0	120.0	360.0	1.0	
4	LP001008	Male	No	0	Graduate	No	6000	0.0	141.0	360.0	1.0	
5	LP001011	Male	Yes	2	Graduate	Yes	5417	4196.0	267.0	360.0	1.0	
...
609	LP002978	Female	No	0	Graduate	No	2900	0.0	71.0	360.0	1.0	
610	LP002979	Male	Yes	3+	Graduate	No	4106	0.0	40.0	180.0	1.0	
611	LP002983	Male	Yes	1	Graduate	No	8072	240.0	253.0	360.0	1.0	
612	LP002984	Male	Yes	2	Graduate	No	7583	0.0	187.0	360.0	1.0	
613	LP002990	Female	No	0	Graduate	Yes	4583	0.0	133.0	360.0	0.0	S

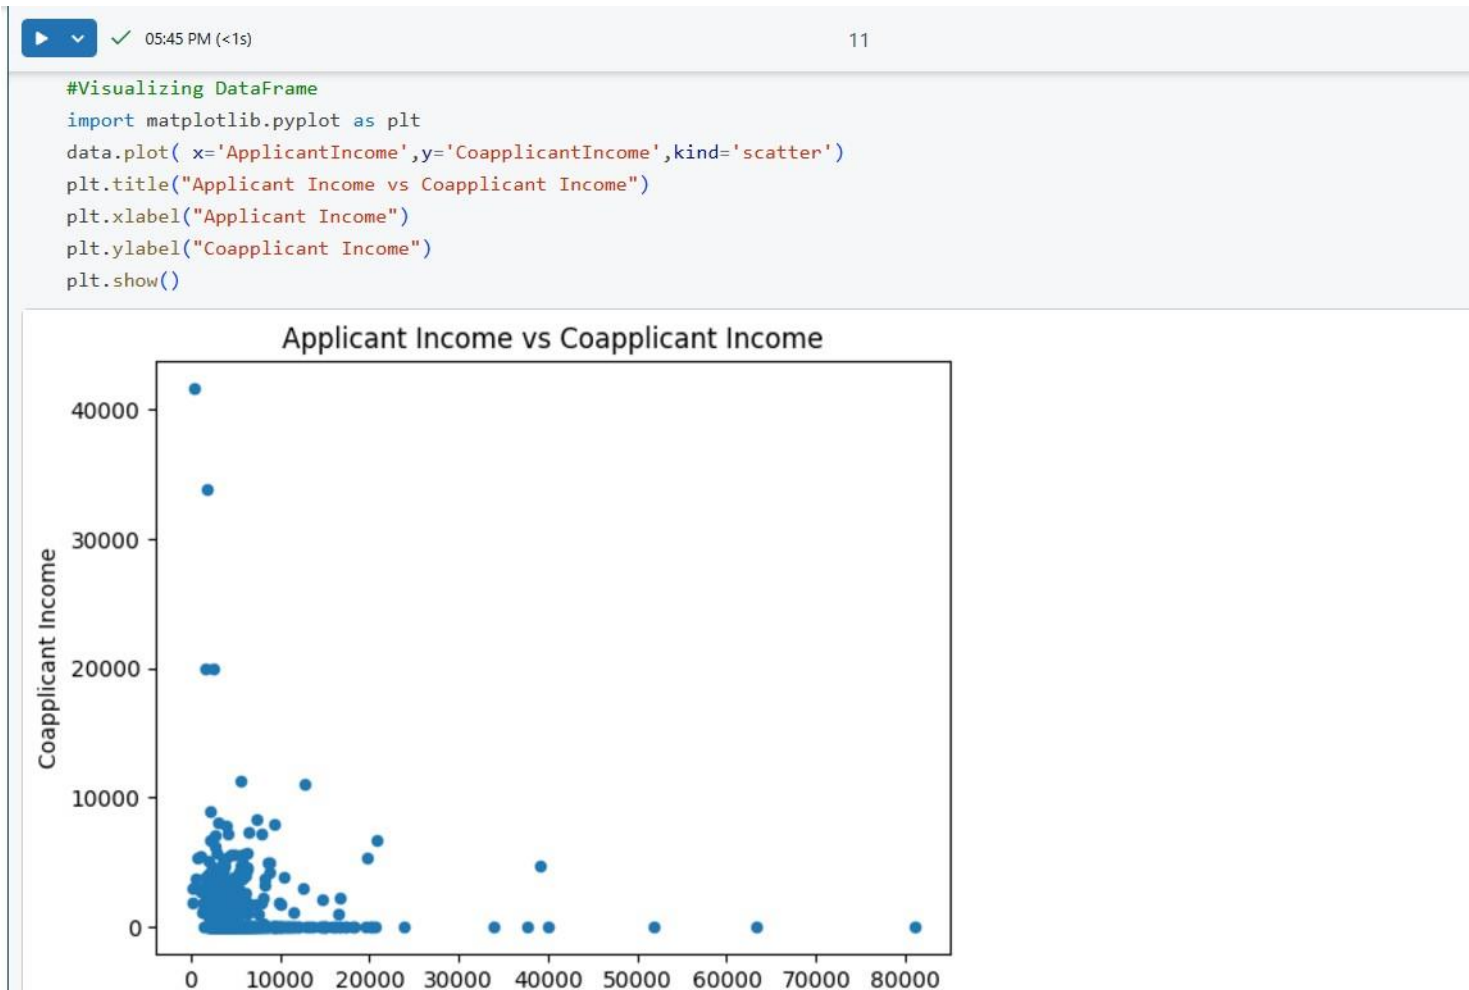
480 rows × 13 columns

```
#Apply Function
```

```
def fun(value):
    if value>3000:
        return 'Yes'
    else:
        return 'No'
```

```
data['newColumn'] = data['ApplicantIncome'].apply(fun)
data.head()
```

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History	Property
0	LP001002	Male	No	0	Graduate	No	5849	0.0	NaN	360.0	1.0	
1	LP001003	Male	Yes	1	Graduate	No	4583	1508.0	128.0	360.0	1.0	
2	LP001005	Male	Yes	0	Graduate	Yes	3000	0.0	66.0	360.0	1.0	
3	LP001006	Male	Yes	0	Not Graduate	No	2583	2358.0	120.0	360.0	1.0	
4	LP001008	Male	No	0	Graduate	No	6000	0.0	141.0	360.0	1.0	



4. Data Exploration and Visualization in Databricks:

05:55 PM (<1s) 1 Python

```
from pyspark.sql import SparkSession

spark = SparkSession.builder.appName("VisualizationExample").getOrCreate()

data = spark.read.format("delta").load("dbfs:/user/hive/warehouse/export")
display(data)
```

(1) Spark Jobs

data: pyspark.sql.dataframe.DataFrame = [id: long, firstName: string ... 6 more fields]

	¹ ₃ id	^A _C firstName	^A _C middleName	^A _C lastName	^A _C gender	^B birthDate	^A _C ssn	¹ ₃ salary
1	1	Pennie	Carry	Hirschmann	F	1955-07-02T04:00:00.000+00:...	981-43-9345	56172
2	2	An	Amira	Cowper	F	1992-02-08T05:00:00.000+00:...	978-97-8086	40203
3	3	Quyen	Marlen	Dome	F	1970-10-11T04:00:00.000+00:...	957-57-8246	53417
4	4	Coralie	Antonina	Marshal	F	1990-04-11T04:00:00.000+00:...	963-39-4885	94727
5	5	Terrie	Wava	Bonar	F	1980-01-16T05:00:00.000+00:...	964-49-8051	79908
6	6	Chassidy	Concepcion	Bourthouloume	F	1990-11-24T05:00:00.000+00:...	954-59-9172	64652
7	7	Geri	Tambra	Mosby	F	1970-12-19T05:00:00.000+00:...	968-16-4020	38195
8	8	Patria	Nancy	Arstall	F	1985-01-02T05:00:00.000+00:...	984-76-3770	102053

```
# Print the schema to understand data types and column structure
data.printSchema()
```

```
root
|-- id: long (nullable = true)
|-- firstName: string (nullable = true)
|-- middleName: string (nullable = true)
|-- lastName: string (nullable = true)
|-- gender: string (nullable = true)
|-- birthDate: timestamp (nullable = true)
|-- ssn: string (nullable = true)
|-- salary: long (nullable = true)
```

```
# Show summary statistics for numeric columns
data.describe().show()
```

▶ (2) Spark Jobs

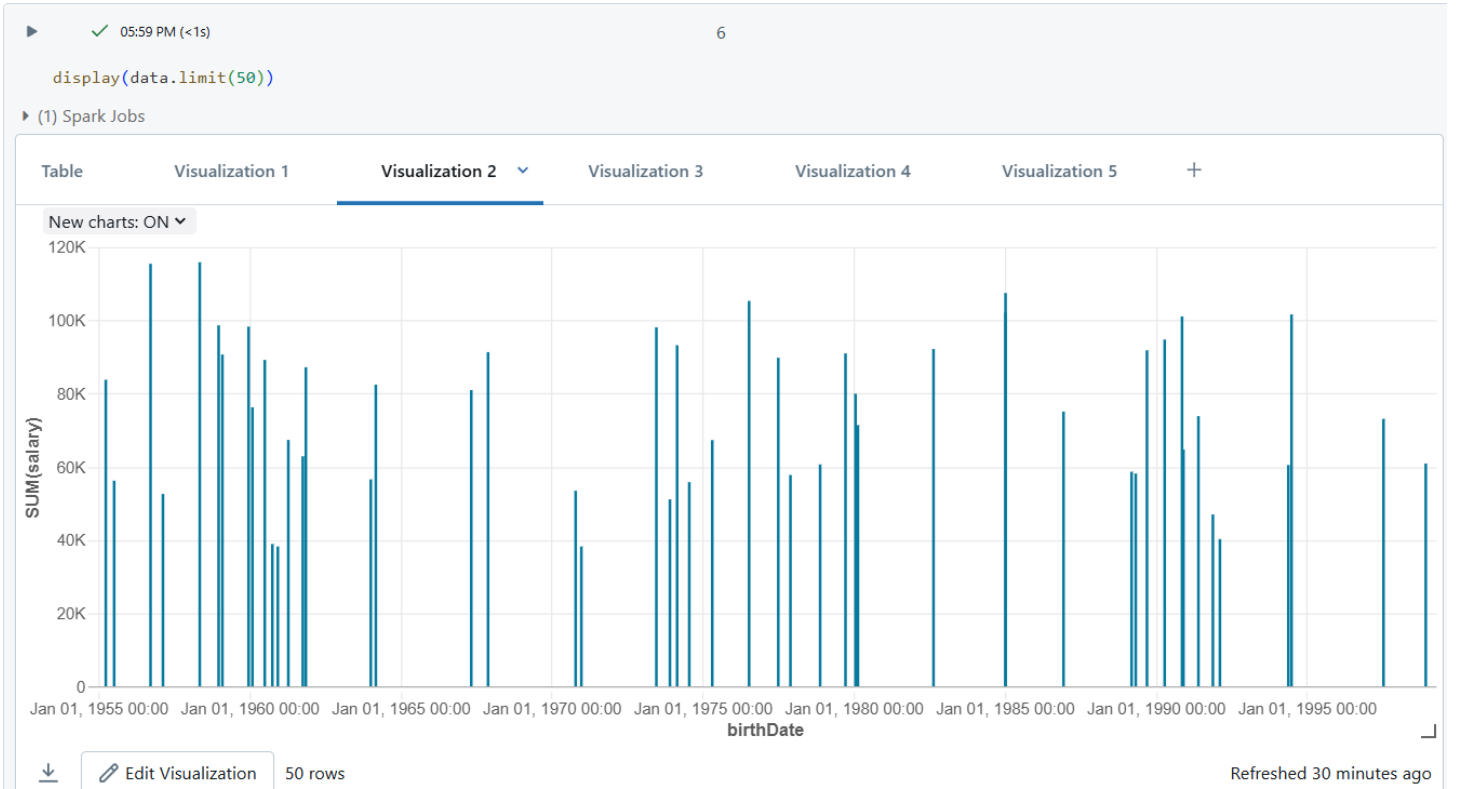

summary	id	firstName	middleName	lastName	gender	ssn	salary
count	1000	1000	1000	1000	1000	1000	1000
mean	500.5	NULL	NULL	NULL	NULL	NULL	72763.54
stddev	288.8194360957494	NULL	NULL	NULL	NULL	NULL	20670.644326853664
min	1	Abbie	Adella	Abraham	F	666-15-8671	-6523
max	1000	Zita	Zulma	Zannutti	F	999-93-2044	134393

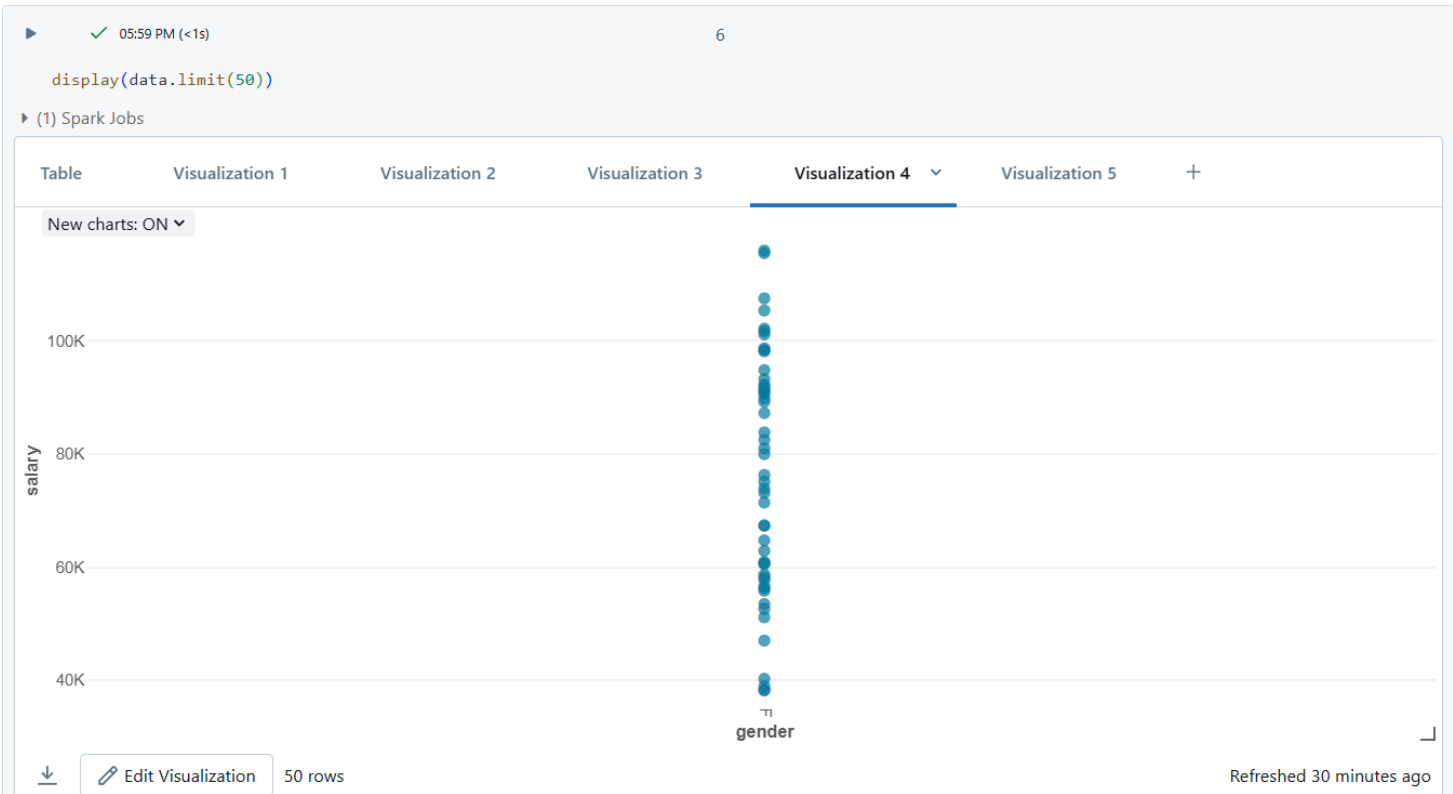
```
# Count the number of rows and columns
num_rows = data.count()
num_cols = len(data.columns)
print(f"Rows: {num_rows}, Columns: {num_cols}")
```

▶ (2) Spark Jobs

Rows: 1000, Columns: 8

Visualization:






```
display(data.limit(50))
```

(1) Spark Jobs

Table

Visualization 1

Visualization 2

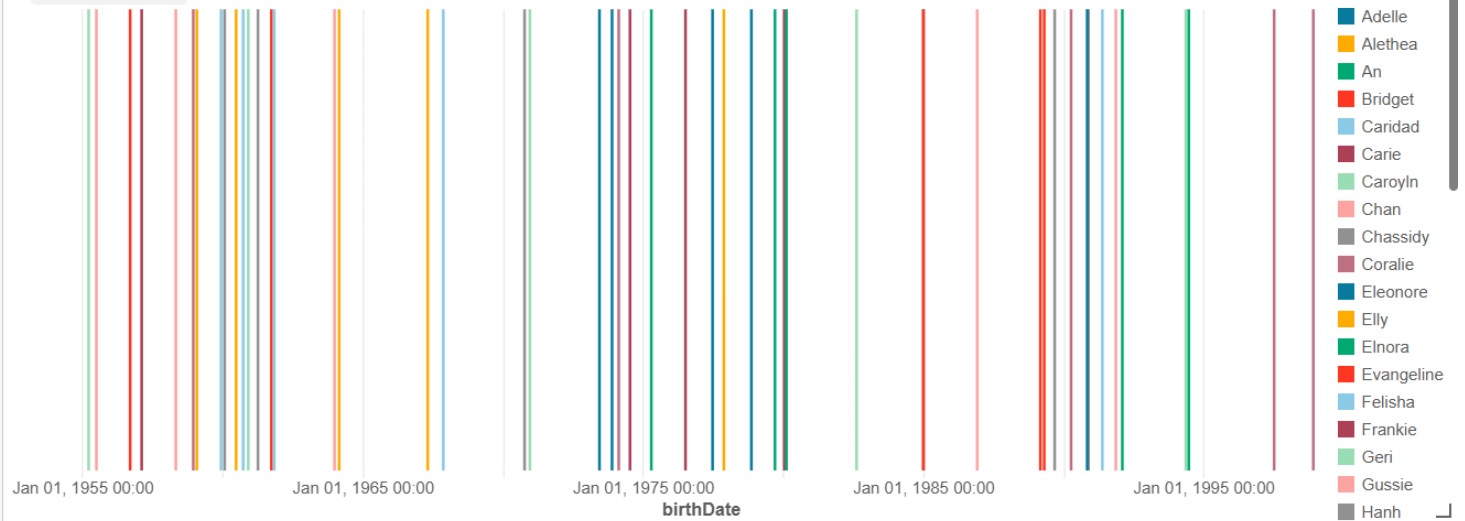
Visualization 3

Visualization 4

Visualization 5

+

New charts: ON



Edit Visualization

50 rows

Refreshed 31 minutes ago